

# Parallel Programming Exercises

Andreas Wilhelm

Technische Universität München

April 20, 2015

# Organization

- ▶ Lecture could start at:
  - ▶ 4:00 PM
  - ▶ 4:15 PM
  - ▶ 4:30 PM
- ▶ Poll is available at <http://doodle.com/ai2kp6872uutyait>
- ▶ Duration: as long as we need, up to 90 min

# Organization

- ▶ Assignments on parallel programming techniques
- ▶ Topics
  - ▶ Pthreads (Posix Threads)
  - ▶ OpenMP (Open Multi-Processing)
  - ▶ Dependency analysis
  - ▶ MPI (Message Passing Interface)
- ▶ Code examples are in C99 (no C++)
- ▶ My email address is: wilhelma@in.tum.de

# Assignments

- ▶ Submission of 80% of the assignments gives 0.3 bonus
- ▶ Submissions will be checked for:
  - ▶ plagiarism
  - ▶ correctness (output, threads, synchronization)
  - ▶ speedup
  - ▶ memory leaks
- ▶ Submission is done on website:  
<https://131.159.74.59/Submission>
  - ▶ requires your LRZ ID and your password
  - ▶ password is not stored and only used for authentication
  - ▶ download SSL-certificate [here](#)
- ▶ Assignment instructions are on the last slides
- ▶ Final exam will contain small programming tasks
- ▶ Solutions will be made public

# Assistance on Assignments

## This week

- ▶ Given by: Andreas Wilhelm
- ▶ Room: 01.04.060
- ▶ Date and Time: Thursday 1:00 PM - 2:30 PM

## After this week

- ▶ Given by: Khalid Alkhalili
- ▶ Email: kfkhalili@hotmail.com
- ▶ Room: 01.04.011
- ▶ Possible Date and Time:
  - ▶ Tuesday 1 PM, 2 PM, 3 PM (90 mins.)
  - ▶ Wednesday 1 PM, 2 PM, 3 PM (90 mins.)
  - ▶ Thursday 1 PM, 2 PM, 3 PM (90 mins.)
- ▶ Poll is available at <http://doodle.com/8mf869zmf5zd8zqc>


# Resources

- ▶ POSIX Threads Programming
- ▶ An Introduction to Parallel Programming, by Peter Pacheco
- ▶ Programming with Posix Threads, by David Butenhof
- ▶ The Linux Programming Interface, by Michael Kerrisk
- ▶ Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill

# Course Prerequisites

- ▶ Knowledge of C
  - ▶ memory management
  - ▶ pointers
  - ▶ global vs. static variables
- ▶ C books
  - ▶ (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
  - ▶ (C99) C Primer Plus, Fifth Edition, by Stephen Prata
- ▶ Experience with Linux Command Line
- ▶ Resources
  - ▶ Book: The Linux Command Line
  - ▶ Basic video introduction: The Shell
- ▶ Knowing GCC
  - ▶ An Introduction to GCC, by Brian Gough

# 50th Anniversary of Moore's Law



Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten  
Dotted line extrapolations by C. Moore



# Year 2005: The Free Lunch Is Over

- ▶ A Fundamental Turn Toward Concurrency in Software
- ▶ Software doesn't get (much) faster with the next microprocessor generation
- ▶ Software developers have to rewrite their applications to use multiple processors in order to speed them up
- ▶ Parallel Programming is hard
  - ▶ to write - complex APIs and needs more code than serial version
  - ▶ to do it correctly - it's easy to introduce new bugs
  - ▶ to debug, order of thread execution is undefined
  - ▶ to make it scalable - will your applications scale with more cores?
  - ▶ better qualified developers are necessary

# Posix Thread Programming

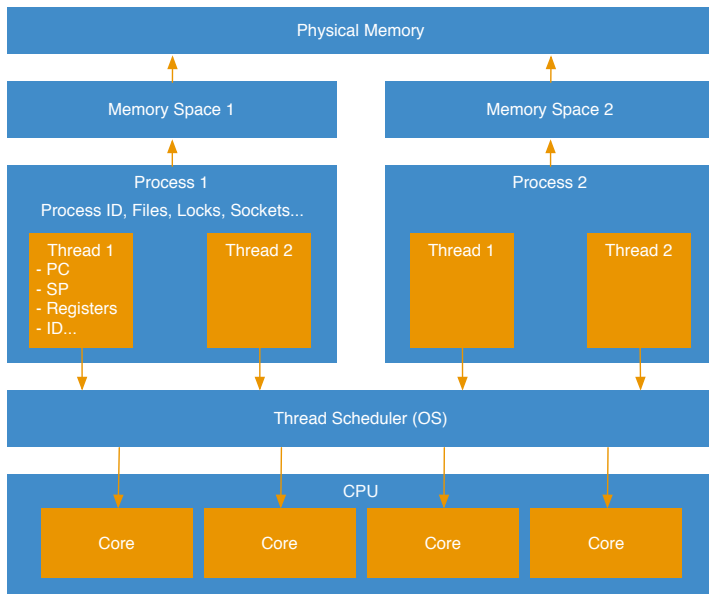
## Definition: Thread

A thread is an independent stream of instructions that can be scheduled to run as such by the operating system.

## POSIX Threads (Pthreads)

- ▶ Were defined in 1995 (IEEE Std 1003.1c-1995)
- ▶ Is an API that defines a set of types, functions and constants
- ▶ Is implemented with a `pthread.h` header and a thread library
- ▶ Functions can be categorized in four groups:
  - ▶ Thread management
  - ▶ Mutexes
  - ▶ Condition variables
  - ▶ Read/write locks and barriers

# Processes vs. Threads




# Why use Multithreading?


- ▶ **Performance gains**  
Parallel processing by multiple processor cores
- ▶ **Increased application throughput**  
Asynchronous system calls possible
- ▶ **Increased application responsiveness**  
Application does not need to block operations
- ▶ **Replacing process-to-process communications**  
Threads may communicate by shared-memory
- ▶ **Efficient use of system resources**  
Lightweight context switches possible
- ▶ **The ability to create well-structured programs**  
Some programs are inherently concurrent

# Concurrency vs. Parallelism

## Concurrency



## Parallelism



# Why are threads 'faster' than processes?

- ▶ Creating a new process with `fork()` has a big overhead: whole memory must be copied
  - ▶ Waste of memory space!
- ▶ Synchronization with processes usually involves system calls.

# Create Pthreads

```
1  int pthread_create(pthread_t *thread,  
2                          const pthread_attr_t *attr,  
3                          void *(*start_routine) (void *),  
4                          void *arg);
```

- ▶ pthread\_t \*thread,
  - ▶ Pointer to thread identifier.
- ▶ const pthread\_attr\_t \*attr
  - ▶ Optional pointer to pthread\_attr\_t to define behavior, if NULL defaults are used.
- ▶ void \*(\*start\_routine) (void \*),
  - ▶ Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- ▶ void \*arg
  - ▶ Pointer to the argument that is used for the executed function.

## Create Pthreads - Example

```
1 void * hello()  
2 {  
3     printf("Hello World from pthread!\n");  
4     return NULL;  
5 }  
6  
7 void main()  
8 {  
9     pthread_t thread;  
10  
11     pthread_create(&thread, NULL, &hello, NULL);  
12  
13     printf("Hello World from main!\n");  
14  
15     pthread_join(thread, NULL);  
16 }
```



# Waiting for Pthread to finish

```
1  int pthread_join(pthread_t thread,  
2                      void **retval);
```

- ▶ pthread\_t thread,
  - ▶ Pointer to thread identifier, for which this function is waiting.
- ▶ void \*\*retval
  - ▶ Optional pointer pointing to a void pointer. This can be used to return data of undefined size.

# Compile & Output

```
gcc --std=gnu99 -pthread -Wall  
    -o hello_world hello_world.c
```

```
Hello World from main!
```

```
Hello World from pthread!
```

# More than One: Hello World with Pthreads Ver. 1

```
1 void main()  
2 {  
3     long num_threads = 3; pthread_t *thread;  
4  
5     thread = malloc((num_threads * sizeof(*thread)));  
6  
7     for (int i = 0; i < num_threads; i++)  
8         pthread_create(thread + i, NULL, &hello, NULL );  
9  
10    for (int i = 0; i < num_threads; i++)  
11        pthread_join(thread[i], NULL );  
12 }
```

# Output

```
[user]$ ./hello_world_thread_1  
Hello World from pthread!  
Hello World from pthread!  
Hello World from pthread!
```

## Single Argument: Hello World with Pthreads Ver. 2

```
1 void * hello(void *ptr)
2 {
3     printf("Hello World from pthread %d!\n", *((int*)ptr));
4     return NULL ;
5 }
```

## Single Argument: Hello World with Pthreads Ver. 2

```
1 void main()
2 {
3     int num_threads = 3;  pthread_t *thread;  int *thread_arg;
4
5     thread = malloc(num_threads * sizeof(*thread));
6     thread_arg = malloc(num_threads * sizeof(*thread_arg));
7
8     for (int i = 0; i < num_threads; i++)
9     {
10         thread_arg[i] = i;
11         pthread_create(thread + i, NULL, &hello, thread_arg + i);
12     }
13
14     for (int i = 0; i < num_threads; i++)
15         pthread_join(thread[i], NULL );
16 }
```

# Output

```
[user]$ ./hello_world_thread_2  
Hello World from pthread 0!  
Hello World from pthread 1!  
Hello World from pthread 2!
```

## Many Arguments: Hello World with Pthreads Ver. 3

```
1  struct pthread_args
2  {
3      long thread_id;
4      long num_threads;
5  };
6
7  void * hello(void *ptr) {
8      struct pthread_args *arg = ptr;
9      printf(
10         "Hello World from pthread %ld of %ld PID = %d TID = %d!\n"
11         arg->thread_id,
12         arg->num_threads,
13         getpid(),
14         (unsigned int)pthread_self());
15
16     return NULL ;
17 }
```



## Many Arguments: Hello World with Pthreads Ver. 3

```
1  void main()
2  {
3      long num_threads = 3;
4      pthread_t *thread;
5      struct pthread_args *thread_arg;
6
7      thread = malloc(num_threads * sizeof(*thread));
8      thread_arg = malloc(num_threads * sizeof(*thread_arg));
9
10     for (int i = 0; i < num_threads; i++)
11     {
12         thread_arg[i].thread_id = i;
13         thread_arg[i].num_threads = num_threads;
14         pthread_create(thread + i, NULL, &hello_pthread, thread_arg + i);
15     }
16
17     for (int i = 0; i < num_threads; i++)
18         pthread_join(thread[i], NULL );
19 }
```

# Output

```
[user]$ ./hello_world_thread_3  
Hello World from pthread 1 of 3 PID = 23750 TID = 23752!  
Hello World from pthread 0 of 3 PID = 23750 TID = 23751!  
Hello World from pthread 2 of 3 PID = 23750 TID = 23753!
```

## Return Result from Pthread in struct Argument

```
1  struct pthread_args
2  {
3      int in, out;
4  };
5
6  void * triple(void *ptr)
7  {
8      struct pthread_args *arg = ptr;
9
10     arg->out = 3 * arg->in;
11
12     return NULL ;
13 }
```

## Return Result from Pthread in struct Argument

```
1 void main() {
2     long num_threads = 3; pthread_t *thread;
3     struct pthread_args *thread_arg;
4
5     thread = malloc(num_threads * sizeof(*thread));
6     thread_arg = malloc(num_threads * sizeof(*thread_arg));
7
8     for (int i = 0; i < num_threads; i++){
9         thread_arg[i].in = i;
10        pthread_create(thread + i, NULL, &triple , thread_arg + i);
11    }
12
13    for (int i = 0; i < num_threads; i++){
14        pthread_join(thread[i], NULL );
15        printf("Triple of %d is %d\n",
16              pthread_args[i].in ,
17              pthread_args[i].out);
18    }
19 }
```

# Return Result from Pthread as Pointer to Memory

```
1
2 void * triple(void *ptr) {
3
4     int *out = malloc(sizeof(*out));
5     *out = 3 * (*(int*)ptr);
6
7     return out;
8 }
```

# Return Result from Pthread as Pointer to Memory

```
1 void main() {
2     long num_threads = 3; pthread_t *thread;
3     int *in;
4
5     thread = malloc(num_threads * sizeof(*thread));
6     in = malloc(num_threads * sizeof(*in));
7
8     for (int i = 0; i < num_threads; i++){
9         in[i] = i;
10        pthread_create(thread + i, NULL, &triple , in + i);
11    }
12
13    for (int i = 0; i < num_threads; i++){
14        int *out;
15        pthread_join(thread[i], &out );
16        printf("Triple of %d is %d\n", in[i], *out);
17        free(out);
18    }
19 }
```

## What have we covered so far?

- ▶ Creating new threads with `pthread_create`
- ▶ Waiting for threads to finish with `pthread_join`
- ▶ Passing arguments to a `pthread` function
- ▶ Returning results from `pthread` function

# Assignment: Histogram

## Task

- ▶ A histogram represents the frequencies of different alphabetic characters in an input file (here: case insensitive).
- ▶ Use POSIX threads to parallelize a given program that computes such a histogram.

## Usage of the program

- ▶ Sequential:  
`./histogram_seq <file> 1 (<#repetitions>)`
- ▶ Parallel:  
`./histogram_par <file> (<#threads>) (<#repet.>)`



## Assignment: Example - "War and Peace" (Tolstoy)

```
1 Process war_and_peace.txt by 1 thread(s)
2
3 315232:      |
4 283708:      |
5 252185:      |
6 220662:      |                |
7 189139: |      |                |                |
8 157616: |      |                | |              | |
9 126092: |      |                | |              | |
10 94569:  |      | |              | |              | |
11 63046:  |      | |              | |              | |
12 31523:  | | | | | | | | | | | | | | | | | | | |
13 0:      | | | | | | | | | | | | | | | | | | | |
14         a b c d e f g h i j k l m n o p q r s t u v w x y z
15
16 Time: 0.01325 seconds
```

## Assignment: Build Histogram (Sequential)

```
1  #include "histogram.h"
2
3  void get_histogram(unsigned int nBlocks,
4                    block_t *blocks,
5                    unsigned int* histogram,
6                    unsigned int num_threads) {
7
8      unsigned int i, j;
9      for (i=0; i<nBlocks; i++) {
10         for (j=0; j<BLOCKSIZE; j++) {
11             if (blocks[i][j] >= 'a' && blocks[i][j] <= 'z')
12                 histogram[blocks[i][j]-'a']++;
13             else if (blocks[i][j] >= 'A' && blocks[i][j] <= 'Z')
14                 histogram[blocks[i][j]-'A']++;
15         }
16     }
17 }
```

# Assignment: Provided Files

- ▶ Makefile
  - ▶ contains rules to build executables
  - ▶ available targets: parallel, sequential, all (default), clean
  - ▶ 'mode=debug make [target]' to build debug version, use 'make clean' before
- ▶ main.c
- ▶ histogram.h
  - ▶ Header file for histogram\_\*.c
- ▶ histogram.c
  - ▶ Defines print\_histogram
- ▶ histogram\_seq.c
  - ▶ Sequential version. Replace it with pi series given in the slides.
- ▶ student/histogram\_par.c
  - ▶ Implement the parallel version in this file
- ▶ war\_and\_peace.txt
  - ▶ Input data: War and peace from Tolstoy

# Assignment: Extract, Build, and Run

## 1. Extract all files to the current directory

```
tar -xvf histogram.tar
```

## 2. Build the program

```
make [sequential] [parallel]
```

- ▶ sequential: build the sequential program
- ▶ parallel: build the parallel program

## 3. Run the sequential program (100 repetitions)

```
student/histogram_seq war_and_peace.txt 1 100
```

## 4. Run the parallel program (with N threads)

```
student/histogram_par war_and_peace.txt N 100
```

# Submission

1. Log into the website
2. Go to Assignments
3. Upload your histogram\_par.c
4. Press Submit

[Parallel Programming](#) [Home](#) [Slides](#) [Assignments](#) [Contact](#)

Welcome ga49qez [Logout](#)

#### NOTE:

1. Some of the gcc versions does not support -Wpedantic compiler flag. If you get an error saying Unrecognized command line option "-Wpedantic", change the flag to -pedantic in Makefile

2. Add -lrt option to LDFlags in case you get Undefined reference to "clock\_gettime" error.

Choose File no file selected

Submit

| Assignment Name    | Submitted By    | Date Of Submission        | Status     |
|--------------------|-----------------|---------------------------|------------|
| Histogram Pthreads | Andreas Wilhelm | April 20, 2015, 6:20 p.m. | ✓ Accepted |

```
Build step successfull
Correctness checks successfull
Pthread checks successfull
Memcheck checks successfull
Helgrind checks successfull
Runtime checks successfull - speedup: 4.1
```